

Latticed-LTL Synthesis in the Presence of Noisy Inputs

Shaul Almagor¹ and Orna Kupferman¹

The Hebrew University, Jerusalem, Israel.

Abstract. In the classical synthesis problem, we are given a linear temporal logic (LTL) formula ψ over sets of input and output signals, and we synthesize a finite-state transducer that realizes ψ : with every sequence of input signals, the transducer associates a sequence of output signals so that the generated computation satisfies ψ . In recent years, researchers consider extensions of the classical Boolean setting to a multi-valued one. We study a setting in which the truth values of the input and output signals are taken from a finite lattice, and the specification formalism is Latticed-LTL (LLTL), where conjunctions and disjunctions correspond to the meet and join operators of the lattice, respectively. The lattice setting arises in practice, for example in specifications involving priorities or in systems with inconsistent viewpoints.

The goal in the latticed setting is to synthesize a transducer that realizes ψ in desired truth values. We solve the LLTL synthesis problem and study cases where the set of attainable truth values are closed under join, thus a maximal attainable value exists, even when the lattice elements are partially ordered.

For the classical synthesis problem, researchers have studied a setting with incomplete information, where the truth values of some of the input signals are hidden and the transducer should nevertheless realize ψ . For the multi-valued setting, we introduce and study a new type of incomplete information, where the truth values of some of the input signals may be noisy, and the transducer should still realize ψ in a desired value. We study the problem of noisy LLTL synthesis, as well as the theoretical aspects of the setting, like the amount of noise a transducer may tolerate, or the effect of perturbing input signals on the satisfaction value of a specification.

1 Introduction

Synthesis is the automated construction of a system from its specification. The basic idea is simple and appealing: instead of developing a system and verifying that it adheres to its specification, we would like to have an automated procedure that, given a specification, constructs a system that is correct by construction. The first formulation of synthesis goes back to Church [10]; the modern approach to synthesis was initiated by Pnueli and Rosner, who introduced LTL (linear temporal logic) synthesis [24]: We are given an LTL formula ψ over sets I and O of input and output signals, and we synthesize a finite-state system that realizes ψ . At each moment in time, the system reads a truth assignment, generated by the environment, to the signals in I , and it generates a truth assignment to the signals in O . Thus, with every sequence of inputs, the transducer associates a sequence of outputs, and it *realizes* ψ if all the computations that are generated by the interaction satisfy ψ . Synthesis has attracted a lot of research and interest [29].

In recent years, researchers have considered extensions of the classical Boolean setting to a multi-valued one, where the atomic propositions are multi-valued, and so is the satisfaction value of specifications. The multi-valued setting arises directly in systems in which the designer can give to the atomic propositions rich values, expressing, for example, energy consumption, waiting time, or different levels of confidence [5, 1], and arises indirectly in probabilistic settings, systems with multiple and inconsistent view-points, specifications with priorities, and

more [21, 14, 2]. Adjusting the synthesis problem to this setting, one works with multi-valued specification formalisms. In such formalisms, a specification ψ maps computations in which the atomic propositions take values from a domain D to a satisfaction value in D . For example, ψ may map a computation in $(\{0, 1, 2, 3\}^{\{p\}})^\omega$ to the maximal value assigned to the atomic proposition p during the computation. Accordingly, the synthesis problem in the multi-valued setting gets as input a specification ψ and a predicate $P \subseteq D$ of desired values, and seeks a system that reads assignments in D^I , responds with assignments in D^O , and generates only computations whose satisfaction value is in P . The synthesis problem has been solved for several multi-valued settings [3, 4, 1].

A different extension of the classical setting of synthesis considers settings in which the system has *incomplete information* about its environment. Early work on incomplete information considers settings in which the system can read only a subset of the signals in I and should still generate only computations that satisfy the specification, which refers to all the signals in $I \cup O$ [16, 25, 19, 6, 7]. The setting corresponds to that of a game with incomplete information, extensively studied in [26]. As shown there, the common practice in handling incomplete information is to move to an exponentially-larger game of complete information, where each state corresponds to a set of states that are indistinguishable by a player with incomplete information in the original game.

More recent work on synthesis with incomplete information studies richer types of incomplete information. In [8], the authors study a setting in which the transducer can read some of the input signals some of the time. In more detail, *sensing* the truth value of an input signal has a cost, the system has a budget for sensing, and it tries to realize the specification while minimizing the required sensing budget. In [31], the authors study games with *errors*. Such games correspond to a synthesis setting in which there are positions during the interaction in which input signals are read by the system with an error. The games are characterized by the number or rate of errors that the system has to cope with, and by the ability of the system to detect whether a current input is erred.

In this work we introduce and study a different model of incomplete information in the multi-valued setting. In our model, the system always reads all input signals, but their value may be perturbed according to a known noise function. This setting naturally models incomplete information in real-life multi-valued settings. For example, when the input is read by sensors that are not accurate (e.g., due to bounded precision, or to probabilistic measuring) or when the input is received over a noisy channel and may come with some distortion. The multi-valued setting we consider is that of *finite lattices*. A lattice is a partially-ordered set $\mathcal{L} = \langle A, \leq \rangle$ in which every two elements ℓ and ℓ' have a least upper bound ($\ell \vee \ell'$) and a greatest lower bound ($\ell \wedge \ell'$). Of special interest are two classes of lattices: (1) Fully ordered, where $\mathcal{L} = \langle \{1, \dots, n\}, \leq \rangle$, for an integer $n \geq 1$ and the usual “less than or equal” order. In this lattice, the operators \vee and \wedge correspond to max and min, respectively. (2) Power-set lattices, where $\mathcal{L} = \langle 2^X, \subseteq \rangle$, for a finite set X , and the containment (partial) order. In this lattice, the operators \vee and \wedge correspond to \cup and \cap , respectively.

The lattice setting is a good starting point to the multi-valued setting. While their finiteness circumvents the infinite-state space of dense multi-values, lattices are sufficiently rich to capture many quantitative settings. Fully-ordered lattices are sometimes useful as is (for example, when modeling priorities [2]), and sometimes thanks to the fact that real values can often be abstracted to finitely many linearly ordered classes. The power-set lattice models a wide range of partially-ordered values. For example, in a setting with inconsistent viewpoints, we have a set X of agents, each with a different viewpoint of the system, and the truth value of a signal or a formula indicates the set of agents according to whose viewpoint the signal or the formula are true. As another example, in a peer-to-peer network, one can refer to the different attributes

of the communication channels by assigning with them subsets of attributes. From a technical point of view, the fact that lattices are partially ordered poses challenges that do not exist in (finite and infinite) full orders. For example, as we are going to see, the fact that a specification is realizable with value ℓ and with value ℓ' does not imply it is realizable with value $\ell \vee \ell'$, which trivially holds for full orders.

We start by defining lattices and the logic *Latticed LTL* (LLTL, for short). We then study theoretical properties of LLTL: the set of attainable satisfaction values of an LLTL formula (e.g., is there a maximal attainable value?), and stability properties, namely the affect of perturbing the values of the atomic propositions on the satisfaction value of formulas. We continue to the synthesis and the noisy synthesis problems for LLTL, which we solve via a translation of LLTL formulas to Boolean automata. We show that by working with universal automata, we can handle the exponential blow-up that incomplete information involves together with the exponential blow-up that determination (or alternation removal, if we take a Safraless approach) involves, thus the synthesis problem stays 2EXPTIME-complete, as it is for LTL.

2 Preliminaries

2.1 Lattices

Consider a set A , a partial order \leq on A , and a subset P of A . An element $\ell \in A$ is an *upper bound* on P if $\ell \geq \ell'$ for all $\ell' \in P$. Dually, ℓ is a *lower bound* on P if $\ell \leq \ell'$ for all $\ell' \in P$. The pair $\langle A, \leq \rangle$ is a *lattice* if for every two elements $\ell, \ell' \in A$, both the least upper bound and the greatest lower bound of $\{\ell, \ell'\}$ exist, in which case they are denoted $\ell \vee \ell'$ (ℓ join ℓ') and $\ell \wedge \ell'$ (ℓ meet ℓ'), respectively. We use $\ell < \ell'$ to indicate that $\ell \leq \ell'$ and $\ell \neq \ell'$. We say that ℓ is a *child* of ℓ' , denoted $\ell \prec \ell'$, if $\ell < \ell'$ and there is no ℓ'' such that $\ell < \ell'' < \ell'$.

A lattice $\mathcal{L} = \langle A, \leq \rangle$ is finite if A is finite. Note that finite lattices are *complete*: every subset of A has a least-upper and a greatest-lower bound. We use \top (*top*) and \perp (*bottom*) to denote the least-upper and greatest-lower bounds of A , respectively. A lattice is *distributive* if for every $\ell_1, \ell_2, \ell_3 \in A$, we have $\ell_1 \wedge (\ell_2 \vee \ell_3) = (\ell_1 \wedge \ell_2) \vee (\ell_1 \wedge \ell_3)$ and $\ell_1 \vee (\ell_2 \wedge \ell_3) = (\ell_1 \vee \ell_2) \wedge (\ell_1 \vee \ell_3)$. The traditional disjunction and conjunction logic operators correspond to the join and meet lattice operators. In a general lattice, however, there is no natural counterpart to negation. A *De-Morgan* (or *quasi-Boolean*) lattice is a lattice in which every element a has a unique complement element $\neg \ell$ such that $\neg \neg \ell = \ell$, De-Morgan rules hold, and $\ell \leq \ell'$ implies $\neg \ell' \leq \neg \ell$. In the rest of this paper we consider only finite distributive De-Morgan lattices. We focus on two classes of such lattices: (1) Fully ordered, where $\mathcal{L} = \langle \{1, \dots, n\}, \leq \rangle$, for an integer $n \geq 1$ and the usual “less than or equal” order. Note that in this lattice, the operators \vee and \wedge correspond to max and min, respectively, and $\neg i = n - i + 1$. (2) Power-set lattices, where $\mathcal{L} = \langle 2^X, \subseteq \rangle$, for a finite set X , and the containment (partial) order. Note that in this lattice, the operators \vee and \wedge correspond to \cup and \cap , respectively, and negation corresponds to complementation.

Consider a lattice $\mathcal{L} = \langle A, \leq \rangle$. A *join irreducible* element in \mathcal{L} is $l \in A$ such that $l \neq \perp$ and for all elements $l_1, l_2 \in A$, if $l_1 \vee l_2 \geq l$, then $l_1 \geq l$ or $l_2 \geq l$. For example, the join irreducible elements in $\langle 2^X, \subseteq \rangle$ are all singletons $\{x\}$, for $x \in X$. By Birkhoff’s representation theorem for finite distributive lattices, in order to prove that $l_1 = l_2$, it is sufficient to prove that for every join irreducible element l it holds that $l_1 \geq l$ iff $l_2 \geq l$. We denote the set of join irreducible elements of \mathcal{L} by $\text{JI}(\mathcal{L})$. For convenience, we often talk about a lattice \mathcal{L} without specifying A and \leq . We then abuse notations and refer to \mathcal{L} as a set of elements and talk about $l \in \mathcal{L}$ or about assignments in \mathcal{L}^{AP} (rather than $l \in A$ or assignments in A^{AP}).

2.2 The logic LLTL

The logic LLTL is a natural generalization of LTL to a multi-valued setting, where the atomic propositions take values from a lattice \mathcal{L} [9, 17]. Given a (finite distributive De-Morgan) lattice \mathcal{L} , the syntax of LLTL is given by the following grammar, where p ranges over a set AP of atomic propositions, and ℓ ranges over \mathcal{L} .

$$\varphi := \ell \mid p \mid \neg\varphi \mid \varphi \vee \varphi \mid \mathbf{X}\varphi \mid \varphi \mathbf{U}\varphi.$$

The semantics of LLTL is defined with respect to a *computation* $\pi = \pi_0, \pi_1, \dots \in (\mathcal{L}^{AP})^\omega$. Thus, in each moment in time the atomic propositions get values from \mathcal{L} . Note that classical LTL coincides with LLTL defined with respect to the two-element fully-ordered lattice. For a position $i \geq 0$, we use π^i to denote the suffix π_i, π_{i+1}, \dots of π . Given a computation π and an LLTL formula φ , the *satisfaction value* of φ in π , denoted $\llbracket \pi, \varphi \rrbracket$, is defined by induction on the structure of φ as follows (the operators on the right-hand side are the join, meet, and complementation operators of \mathcal{L}).

- $\llbracket \pi, \ell \rrbracket = \ell$.
- $\llbracket \pi, \varphi \vee \psi \rrbracket = \llbracket \pi, \varphi \rrbracket \vee \llbracket \pi, \psi \rrbracket$.
- $\llbracket \pi, p \rrbracket = \pi_0(p)$.
- $\llbracket \pi, \neg\varphi \rrbracket = \neg \llbracket \pi, \varphi \rrbracket$.
- $\llbracket \pi, \mathbf{X}\varphi \rrbracket = \llbracket \pi^1, \varphi \rrbracket$.
- $\llbracket \pi, \varphi \mathbf{U}\psi \rrbracket = \bigvee_{i \geq 0} (\llbracket \pi^i, \psi \rrbracket \wedge \bigwedge_{0 \leq j < i} \llbracket \pi^j, \varphi \rrbracket)$.

Example 1. Consider a setting in which three agents a, b , and c have different view-points on a system \mathcal{S} . A truth assignment for the atomic propositions is then a function in $(2^{\{a,b,c\}})^{AP}$ assigning to each $p \in AP$ the set of agents according to whose view-point p is true. We reason about \mathcal{S} using the lattice $\mathcal{L} = \langle 2^{\{a,b,c\}}, \subseteq \rangle$. For example, the truth value of the formula $\psi = G(\text{req} \rightarrow F \text{grant})$ in a computation is the set of agents according to whose view-point, whenever a request is sent, it is eventually granted.

Remark 1. Recall that the constants **True** and **False** in LTL do not add to its expressive power. Indeed, **True** can be replaced by $p \vee (\neg p)$, for a Boolean atomic proposition p , and similarly for **False**. This is not the case for the constants $\ell \in \mathcal{L}$ in LLTL. In particular, constants can be used to upper or lower bound the satisfaction value of an LLTL formula. For example, the truth value of the LLTL formula $\{a, b\} \wedge \psi$, defined with respect to the lattice $\langle 2^{\{a,b,c\}}, \subseteq \rangle$, is the set of agents that is both a subset of $\{a, b\}$ and according to whose viewpoint, the specification ψ is satisfied.

2.3 LLTL synthesis

Consider a lattice \mathcal{L} and finite disjoint sets I and O of input and output signals that take values in \mathcal{L} . An (I/O) -transducer over \mathcal{L} models an interaction between an environment that generates in each moment in time an input in \mathcal{L}^I and a system that responds with outputs in \mathcal{L}^O . Formally, an (I/O) -transducer over \mathcal{L} (transducer, when I, O , and \mathcal{L} are clear from the context) is a tuple $\mathcal{T} = \langle \mathcal{L}, I, O, S, s_0, \eta, \tau \rangle$ where S is a finite set of states, $s_0 \in S$ is an initial state, $\eta : S \times \mathcal{L}^I \rightarrow S$ is a deterministic transition function, and $\tau : S \rightarrow \mathcal{L}^O$ is a labeling function. We extend η to words in $(\mathcal{L}^I)^*$ in the straightforward way. Thus, $\eta : (\mathcal{L}^I)^* \rightarrow S$ is such that $\eta(\epsilon) = s_0$, and for $x \in (\mathcal{L}^I)^*$ and $i \in \mathcal{L}^I$ we have $\eta(x \cdot i) = \eta(\eta(x), i)$. Each transducer \mathcal{T} induces a *strategy* $f_{\mathcal{T}} : (\mathcal{L}^I)^* \rightarrow \mathcal{L}^O$ where for all $w \in (\mathcal{L}^I)^*$ we have $f_{\mathcal{T}}(w) = \tau(\eta(w))$. Thus, $f_{\mathcal{T}}(w)$ is the letter that \mathcal{T} outputs after reading the sequence w of input letters.

Given a sequence $i_0, i_1, i_2, \dots \in (\mathcal{L}^I)^\omega$ of input assignments, the transducer generates the computation $\rho = (i_0 \cup o_0), (i_1 \cup o_1), (i_2 \cup o_2), \dots \in (\mathcal{L}^{I \cup O})^\omega$, where for all $j \geq 1$ we have $o_j = f_{\mathcal{T}}(i_0 \dots i_{j-1})$.

Consider a lattice \mathcal{L} , an LLTL formula φ over the atomic propositions $I \cup O$, and a predicate $P \subseteq \mathcal{L}$. We say that a transducer \mathcal{T} *realizes* $\langle \varphi, P \rangle$ if for every computation ρ of \mathcal{T} , it holds that $\llbracket \rho, \varphi \rrbracket \in P$. The *realizability* problem for LLTL is to determine, given φ and P , whether there exists a transducer that realizes $\langle \varphi, P \rangle$. We then say that φ is (I/O) -realizable with values in P . The *synthesis* problem is then to generate such a transducer. Of special interest are predicates P that are upward closed. Thus, there is a set $X \subseteq \mathcal{L}$ such that for all $\ell \in \mathcal{L}$ we have that $\ell \in P$ iff there is $\ell' \in X$ such that $\ell \geq \ell'$.

2.4 Noisy synthesis

Consider an LLTL formula φ over atomic proposition $I \cup O$ and a predicate P . In *noisy synthesis*, we consider the synthesis problem in a setting in which the inputs are read with some perturbation and the goal is to synthesize a transducer that nevertheless realizes $\langle \varphi, P \rangle$.

In order to formalize the above intuition, we first formalize the notion of noise. Consider a lattice $\mathcal{L} = \langle A, \leq \rangle$ and two elements $\ell_1, \ell_2 \in \mathcal{L}$. We define the *distance* between ℓ_1 and ℓ_2 , denoted $d(\ell_1, \ell_2)$, as the shortest path from ℓ_1 to ℓ_2 in the undirected graph $\langle A, E_{\prec} \rangle$ in which $E_{\prec}(v, v')$ iff $v \prec v'$ or $v' \prec v$. For example, in the fully-ordered lattice \mathcal{L} , we have $d(i, j) = |i - j|$, and in the power-set lattice, the distance coincides with the Hamming distance, thus $d(X_1, X_2) = |(X_1 \setminus X_2) \cup (X_2 \setminus X_1)|$. For two assignments $f, f' \in \mathcal{L}^{AP}$, we define $d(f, f') = \max_{p \in AP} d(f(p), f'(p))$.

We assume we are given a *noise function* $\rho : \mathcal{L}^I \rightarrow 2^{\mathcal{L}^I}$, describing the possible perturbations of each input. That is, for every $i \in \mathcal{L}^I$, the set $\rho(i)$ consists of the inputs that the system may read when the environment generates i . A natural noise function is $\rho(i) = \{j : d(i, j) \leq \gamma\}$, for some constant γ , which is the γ -units ball around i . Given a noise function ρ and two computations $\pi, \pi' \in (\mathcal{L}^{I \cup O})^\omega$, we say that π' is ρ -indistinguishable from π if for every $i \geq 0$ we have that $\pi'_i|_I \in \rho(\pi_i|_I)$ and $\pi'_i|_O = \pi_i|_O$, where $\sigma|_I$ is the restriction of $\sigma \in \mathcal{L}^{I \cup O}$ to inputs in I , and similarly for O . Thus, π' is obtained from π by changing only the assignment to input signals, within ρ . Note that ρ need not be a symmetric function, nor is the definition of ρ -indistinguishability. We say that a transducer \mathcal{T} *realizes* $\langle \varphi, P \rangle$ *with noise* ρ if for every computation π of \mathcal{T} we have that $\llbracket \pi', \varphi \rrbracket \in P$ for all computations π' that are ρ -indistinguishable from π . Thus, the reaction of \mathcal{T} on every input sequence satisfies φ in a desired satisfaction value even if the input sequence is read with noise ρ .

2.5 Automata and games

An *automaton over infinite words* is $\mathcal{A} = \langle \Sigma, Q, Q_0, \delta, \alpha \rangle$, where Σ is the input alphabet, Q is a finite set of states, $Q_0 \subseteq Q$ is a set of initial states, $\delta : Q \times \Sigma \rightarrow 2^Q$ is a transition function, and α is an acceptance condition. When \mathcal{A} is a *generalized Büchi* or a *generalized co-Büchi* automaton, then $\alpha \subseteq 2^Q$ is a set of sets of accepting states. When \mathcal{A} is a *parity* automaton, then $\alpha = (F_1, \dots, F_d)$ where the sets in α form a partition of Q . The number of sets in α is the *index* of \mathcal{A} . An automaton is *deterministic* if $|Q_0| = 1$ and for every $q \in Q, \sigma \in \Sigma$ we have that $|\delta(q, \sigma)| = 1$. A run $r = r_0, r_1, \dots$ of \mathcal{A} on a word $w = w_1 \cdot w_2 \dots \in \Sigma^\omega$ is an infinite sequence of states such that $r_0 \in Q_0$, and for every $i \geq 0$, we have that $r_{i+1} \in \delta(r_i, w_{i+1})$. We denote by $\text{inf}(r)$ the set of states that r visits infinitely often, that is $\text{inf}(r) = \{q : r_i = q \text{ for infinitely many } i \in \mathbb{N}\}$. The run r is *accepting* if it satisfies α . For generalized Büchi automata, the run has to visit all the sets in α infinitely often. Formally, for

every set $F \in \alpha$ we have that $\inf(r) \cap F \neq \emptyset$. Dually, in generalized co-Büchi automata, there should exist a set $F \in \alpha$ for which $\inf(r) \cap F = \emptyset$. For parity automata, the run r is accepting if the minimal index i for which $\inf(r) \cap F_i \neq \emptyset$ is even.

When \mathcal{A} is a *nondeterministic* automaton, it accepts a word w if it has an accepting run of on w . When \mathcal{A} is a *universal* automaton, it accepts a word w if all its runs on w are accepting. The language of \mathcal{A} , denoted $L(\mathcal{A})$, is the set of words that \mathcal{A} accepts. Note that a deterministic automaton has a single run of each word and can thus be viewed as both a nondeterministic and a universal automaton.

A *parity game* is $\mathcal{G} = \langle \Sigma_1, \Sigma_2, S, s_0, \delta, \alpha \rangle$, where Σ_1 and Σ_2 are alphabets for Players 1 and 2, respectively, S is a finite set of states, $s_0 \in S$ is an initial state, $\delta : S \times \Sigma_1 \times \Sigma_2 \rightarrow S$ is a transition function, and $\alpha = \{F_1, \dots, F_d\}$ is a parity acceptance condition, as described above. A play of the game starts in s_0 . In each turn Player 1 chooses a letter $\sigma \in \Sigma_1$ and Player 2 chooses a letter $\tau \in \Sigma_2$. The play then moves from the current state s to the state $\delta(s, \sigma, \tau)$. Formally, a *play* of \mathcal{G} is an infinite sequence $\rho = \langle s_0, \sigma_0, \tau_0 \rangle, \langle s_1, \sigma_1, \tau_1 \rangle, \dots$ such that for every $i \geq 0$, we have that $s_{i+1} = \delta(s_i, \sigma_i, \tau_i)$. We define $\inf(\rho) = \{s \in S : s = s_i \text{ for infinitely many } i \in \mathbb{N}\}$. A play ρ is *winning for Player 1* if the minimal index i for which $\inf(\rho) \cap F_i \neq \emptyset$ is even. A *strategy* for Player 1 is a function $f : (S \times \Sigma_1 \times \Sigma_2)^* \times S \rightarrow \Sigma_1$ that assigns, for every finite prefix of a play, the next move for Player 1. Similarly, a strategy for Player 2 is a function $g : (S \times \Sigma_1 \times \Sigma_2)^* \times S \times \Sigma_1 \rightarrow \Sigma_2$. A pair of strategies f, g for Players 1 and 2, respectively, induces a single play that conforms with the strategies. A strategy is *memoryless* if it does not depend on the history of the play. Thus, a memoryless strategy for Player 1 is a function $f : S \rightarrow \Sigma_1$ and for Player 2 it is a function $g : S \times \Sigma_1 \rightarrow \Sigma_2$.

We say that Player 1 wins \mathcal{G} if there exists a strategy f for Player 1 such that for every strategy g for Player 2, the play induced by f and g is winning for Player 1. We say that Player 2 wins \mathcal{G} if there exists a strategy g for Player 2 such that for every strategy f of Player 1, the play induced by f and g is not winning for Player 1.

2.6 Solving the Boolean synthesis problem

The classical solution for the synthesis problem for LTL goes via games [24].¹ It involves a translation of the specification into a deterministic parity word automaton (DPW) over the alphabet $2^{I \cup O}$, which is then transformed into a game in which the players alphabets are 2^I and 2^O . More recent solutions avoid the determination and the solution of parity games and use instead alternating tree automata [20, 13]. The complexity of both approaches coincide. Below we describe the classical solution for the synthesis problem, along with its complexity, when the starting point is a specification given by a DPW.² In Remark 2, we describe an alternative, Safraless, approach, where the starting point is a universal co-Büchi automaton.

Theorem 1. *Consider a specification φ over I and O given by means of a DPW \mathcal{D}_φ over the alphabet $2^{I \cup O}$, with n states and index k . The synthesis problem for φ can be solved in time $O(n^k)$.*

Proof. Let $\mathcal{D}_\varphi = \langle 2^{I \cup O}, Q, q_0, \delta, \alpha \rangle$. Let \mathcal{G}_φ be a game that models an interaction that simulates \mathcal{D}_φ between a system (Player 1) that generates assignments in 2^O and an environment (Player 2) that generates assignments in 2^I . Formally, $\mathcal{G}_\varphi = \langle 2^O, 2^I, Q, q_0, \eta, \alpha \rangle$, where $\eta : Q \times 2^O \times 2^I$ is such that for every $q \in Q$, $i \in 2^I$, and $o \in 2^O$, we have that $\eta(q, i, o) =$

¹ In [24] and other early works the games are formulated by means of tree automata.

² State of the art algorithms for solving parity games achieve a better complexity [15, 28]. The bound, however, remains polynomial in n and exponential in k . Since the challenge of solving parity games is orthogonal to our contribution here, we keep this component of our contribution simple.

$\delta(q, i \cup o)$. By [11], the game is determined and one of the players has a memoryless winning strategy. A winning memoryless strategy for Player 1 in \mathcal{G}_φ is then a transducer that realizes φ . The game \mathcal{G}_φ has n states and index k . Hence, by [15, 28], we can find a memoryless strategy for the winner in time $O(n^k)$. \square

3 Properties of LLTL

In this section we study properties of the logic LLTL. We focus on the set of attainable satisfaction values of an LLTL formula and on stability properties, namely the affect of perturbing the values of the atomic propositions on the satisfaction value of formulas.

3.1 Attainable values

Consider a lattice \mathcal{L} . We say that \mathcal{L} is *pointed* if for all LLTL formulas φ , partitions $I \cup O$ of AP , and values $\ell_1, \ell_2 \in \mathcal{L}$, if φ is (I/O) -realizable with value ℓ_1 and with value ℓ_2 , then φ is also (I/O) -realizable with value $\ell_1 \vee \ell_2$. Observe that if \mathcal{L} is pointed, then every LLTL formula over \mathcal{L} has a single transducer that realizes it with a maximal value.

We start by showing that in general, not all lattices are pointed. In fact, our example has $O = \emptyset$, where (I/O) -realizability coincides with satisfiability.

Theorem 2. *Not all distributive De-Morgan lattices are pointed.*

Proof. Consider the lattice $\mathcal{L} = \langle 2^{\{a,b\}} \times \{0, 1\}, \leq \rangle$ where $\langle S_1, v_1 \rangle \leq \langle S_2, v_2 \rangle$ iff $v_1 \leq v_2$ or $(v_1 = v_2 \text{ and } S_1 \subseteq S_2)$. (See Figure 1). We define $\neg \langle S, v \rangle = \langle \{a, b\} \setminus S, 1 - v \rangle$. That is, negation negates both components. It is easy to verify that \mathcal{L} is a distributive De-Morgan lattice.

Let $I = \{p\}$ and consider the formula $\varphi = (p \wedge \langle \{a\}, 1 \rangle) \vee (\neg p \wedge \langle \{b\}, 1 \rangle)$. Both $\langle \{a\}, 1 \rangle$ and $\langle \{b\}, 1 \rangle$ are attainable satisfaction values of φ . For example, by setting p to $\langle \{a\}, 1 \rangle$ or to $\langle \{a\}, 0 \rangle$. On the other hand, for every assignment ℓ to p , the second component of either ℓ or $\neg \ell$ is 0. Consequently, $\langle \{a, b\}, 1 \rangle$ is not attainable, thus \mathcal{L} is not pointed. \square




Fig. 1. The lattice $\langle 2^{\{a,b\}} \times \{0, 1\}, \leq \rangle$




Fig. 2. An N5 structure

Theorem 3. *Fully-ordered lattices and power-set lattices are pointed.*

Proof. For fully-ordered lattices, we have $\ell_1 \vee \ell_2 \in \{\ell_1, \ell_2\}$, so pointed-ness is obvious. We prove the claim for power-set lattices. Consider a lattice $\mathcal{L} = \langle 2^X, \subseteq \rangle$ for some finite set X , and consider an LLTL formula φ over the atomic propositions $I \cup O$. For every set $\ell \in \mathcal{L}$ and element $x \in X$, we define the *projection* $\ell|_x$ of ℓ on x to be **True** if $x \in \ell$ and **False** otherwise. We extend the definition of projection to a letter $\sigma \in \mathcal{L}^{I \cup O}$ by letting $p \in \sigma|_x$

iff $\mathcal{L}(p)_{|x} = \text{True}$. Thus, $\sigma_{|x} \subseteq I \cup O$. Finally, we extend the definition to a computation $\pi \in (\mathcal{L}^{I \cup O})^\omega$ by setting $(\pi_{|x})_i = (\pi_i)_{|x}$. Observe that $\pi_{|x} \in (2^{I \cup O})^\omega$.

For an element $x \in X$, let $\varphi_{|x}$ be the LTL formula obtained from φ by replacing every element $\ell \in \mathcal{L}$ that appears in φ by $\ell_{|x}$. Since the syntax of LLTL differs from that of LTL only by allowing elements from the lattice, it follows that $\varphi_{|x}$ is indeed an LTL formula.

We prove that for every computation $\pi \in (\mathcal{L}^{I \cup O})^\omega$ and for every $\ell \in \mathcal{L}$, it holds that $\llbracket \pi, \varphi \rrbracket \geq \ell$ iff $\pi_{|x} \models \varphi_x$ for all $x \in \ell$. Observe that $\llbracket \pi, \varphi \rrbracket \geq \ell$ iff $x \in \llbracket \pi, \varphi \rrbracket$ for all $x \in \ell$. From here the claim easily follows by induction on the structure of φ .

Now, assume that φ is (I/O) -realizable with value at least ℓ_1 and with value at least ℓ_2 . We claim that φ is realizable with value $\ell_1 \vee \ell_2$. W.l.o.g we can assume $\ell_1 \cap \ell_2 = \emptyset$ (otherwise we replace ℓ_2 with $\ell_2 \setminus \ell_1$). Consider an environment-computation $\tau \in (\mathcal{L}^I)^\omega$. Since φ is realizable with value at least ℓ_1 and realizable with value at least ℓ_2 , there exist computations $\rho, \rho' \in \mathcal{L}^{I \cup O}$ both of which have the input computation τ , such that $\llbracket \rho, \varphi \rrbracket \geq \ell_1$ and $\llbracket \rho', \varphi \rrbracket \geq \ell_2$. We construct an output-computation $\eta \in \mathcal{L}^O$ that such that the combination of η with τ satisfies φ with value at least $\ell_1 \vee \ell_2$ as follows. For every index $i \geq 0$ and $o \in O$, we define $\eta_i(o) = (\ell_1 \cap \rho_i(o)) \cup (\ell_2 \cap \rho'_i(o))$.

Consider the computation π obtained by combining τ and η . For every $x \in \ell_1$, we have that $\pi_{|x} = \rho_{|x}$. Thus, $\pi_{|x} \models \varphi_x$ for every $x \in \ell_1$. Similarly, for every $x \in \ell_2$, we have that $\pi_{|x} = \rho'_{|x}$, so $\pi_{|x} \models \varphi_x$ for every $x \in \ell_2$. We conclude that for every $x \in \ell_1 \cup \ell_2$ it holds that $\pi_{|x} \models \varphi$, and so $\llbracket \pi, \varphi \rrbracket \geq \ell_1 \vee \ell_2$. \square

3.2 Stability

For two computations $\pi = \pi_0, \pi_1, \dots$ and $\pi' = \pi'_0, \pi'_1, \dots$, both in $(\mathcal{L}^{AP})^\omega$, we define the *global distance* between π and π' , denoted $gd(\pi, \pi')$, as $\sum_{i \geq 0} d(\pi_i, \pi'_i)$. Note that $gd(\pi, \pi')$ may be infinite. We define the *local distance* between π and π' , denoted $ld(\pi, \pi')$, as $\max_{i \geq 0} d(\pi_i, \pi'_i)$. Note that $ld(\pi, \pi') \leq |\mathcal{L}|$.

Consider an LLTL formula φ over AP and \mathcal{L} . We say that φ is *globally stable* if for every pair π and π' of computations, we have $d(\llbracket \pi, \varphi \rrbracket, \llbracket \pi', \varphi \rrbracket) \leq gd(\pi, \pi')$. Thus, the difference between the satisfaction value of φ in π and π' is bounded by the sum of differences between matching locations in π and π' . Also, φ is *locally stable* if for every pair π and π' of computations, we have $d(\llbracket \pi, \varphi \rrbracket, \llbracket \pi', \varphi \rrbracket) \leq ld(\pi, \pi')$. Thus, the difference between the satisfaction value of φ in π and π' is bounded by the maximal difference between matching locations in π and π' . Here, we study stability of all LLTL formulas. In Section 5.3, we study the problem of deciding whether a given LLTL formula is stable, and discuss the relevancy of stability to synthesis with noise.

Consider an LLTL formula φ over the atomic propositions AP , and consider computations $\pi, \pi' \in (\mathcal{L}^{AP})^\omega$. Assume that $gd(\pi, \pi') \leq 1$. That is, π and π' differ only in one location, where they differ in the value of a single atomic proposition, whose value in π is a child of its value in π' or vice versa. It is tempting to think that then, $d(\llbracket \pi, \varphi \rrbracket - \llbracket \pi', \varphi \rrbracket) \leq 1$. That is, φ should be globally stable. We start by breaking this intuition, showing that for non-distributive lattices, this is false.

Theorem 4. *LLTL formulas may not be globally stable with respect to non-distributive lattices.*

Proof. Consider the lattice N5 depicted in Figure 2. Consider the formula $\varphi = p \vee q$, and a computation π such that $\pi_0(p) = s$ and $\pi_0(q) = x$. Clearly $\llbracket \pi, \varphi \rrbracket = x$. Now, let π' be the computation obtained from π by setting $\pi'_0(p) = z$. It holds that $gd(\pi, \pi') = 1$. However, $\llbracket \pi', \varphi \rrbracket = z \vee x = w$, and $d(x, w) = 2$. Thus, φ is not globally stable over the lattice N5. \square

In a distributive lattice \mathcal{L} , the structure of N5 is never a sub-lattice. Formally, an *N5 structure* in a lattice \mathcal{L} is a tuple $\langle x, y, z, w, s \rangle$ such that the following relations hold (See Figure 2): $s < x < y < w$, $s < z < w$, $y \not\leq z$, $z \not\leq y$, $x \not\leq z$, and $z \not\leq x$. Note that $x \vee (z \wedge y) = x \vee s = x$, whereas $(x \vee z) \wedge (x \vee y) = w \wedge y = y$. Hence, a distributive lattice \mathcal{L} does not contain elements $\langle x, y, z, w, s \rangle$ that form an N5 structure. As we now show, this can be used in order to prove that when defined with respect to a distributive lattice, all LLTL formulas are globally stable.

Theorem 5. *LLTL formulas are globally stable with respect to De-Morgan distributive lattices.*

Proof. We prove that for every LLTL φ and computations $\pi, \pi' \in (\mathcal{L}^{AP})^\omega$, if $d(\pi, \pi') = 1$, then $d(\llbracket \pi, \varphi \rrbracket, \llbracket \pi', \varphi \rrbracket) \leq 1$. The result then follows by induction on $d(\pi, \pi')$.

Consider an LLTL formula φ and computations $\pi, \pi' \in (\mathcal{L}^{AP})^\omega$ such that $d(\pi, \pi') = 1$. That is, there exists a single index $i \geq 0$ such that $d(\pi_i, \pi'_i) = 1$ and $\pi_j = \pi'_j$ for all $j \neq i$. W.l.o.g, there is $p \in AP$ such that $\pi_i(p) \preceq \pi'_i(p)$. By Birkhoff's representation theorem, there exists a unique element $u \in \text{JI}(\mathcal{L})$ such that $\pi'_i(p) = \pi_i(p) \vee u$.

We prove, by induction over the structure of φ , that $\llbracket \pi, \varphi \rrbracket \in \{\llbracket \pi, \varphi \rrbracket \wedge \neg u, \llbracket \pi, \varphi \rrbracket, \llbracket \pi, \varphi \rrbracket \vee u\}$ and that $d(\llbracket \pi', \varphi \rrbracket, \llbracket \pi, \varphi \rrbracket) \leq 1$.

The proof is detailed in Appendix A.1. As detailed there, the interesting case is when $\varphi = \psi \vee \theta$, where we use the fact that a distributed lattice cannot have an N5 structure. \square

We now turn to study local stability. Since local stability refers to the maximal change along a computation, it is a very permissive notion. In particular, it is not hard to see that in a fully-ordered lattice, a local change of 1 entails a change of at most 1 in the satisfaction value. Thus, we have the following.

Theorem 6. *LLTL formulas are locally stable with respect to fully-ordered lattices.*

In partially-ordered lattices, however, things are more involved, as local changes may be in different “directions”. Formally, we have the following.

Theorem 7. *LLTL formulas may not be locally stable.*

Proof. Consider the power-set lattice $\langle 2^{a,b}, \subseteq \rangle$ and the LLTL formula $\varphi = p \vee \times p$. Consider computations π and π' with $\pi_0(p) = \pi_1(p) = \emptyset$, $\pi'_0(p) = \{a\}$, and $\pi'_1(p) = \{b\}$. It holds that $ld(\pi, \pi') = 1$, whereas $d(\llbracket \pi, \varphi \rrbracket, \llbracket \pi', \varphi \rrbracket) = d(\emptyset, \{a, b\}) = 2$. We conclude that φ is not locally stable. \square

4 Translating LLTL to Automata

In this section we describe an automata-theoretic approach for reasoning about LLTL specifications. One approach is to develop a framework that is based on *lattice automata* [17]. Like LLTL formulas, lattice automata map words to values in a lattice. Lattice automata have proven to be useful in solving the satisfiability and the model-checking problems for LLTL [17]. However, the solution of the synthesis problem involves automata-theoretic constructions for which the latticed counterpart is either not known or is very complicated. In particular, Safra's determinization construction has not yet been studied for lattice automata, and a latticed counterpart of it is not going to be of much fun. Likewise, the solution of two-player games (even reachability, and moreover parity) in the latticed setting is much more complicated than in the Boolean setting. In particular, obtaining a value $\ell_1 \vee \ell_2$ in a latticed game may require one strategy for obtaining ℓ_1 and a different strategy for obtaining ℓ_2 [18]. When the game is induced by

a realizability problem, it is not clear how to combine such strategies into a single transducer that realizes the underlying specification with value $\ell_1 \vee \ell_2$.

Accordingly, a second approach, which is the one we follow, is to use Boolean automata. The fact that LLTL formulas have finitely many possible satisfaction values suggests that this is possible. For fully-ordered lattices, a similar approach has been taken in [12, 1]. Beyond the challenge in these works of maintaining the simplicity of the automata-theoretic framework of LTL, an extra challenge in the latticed setting is caused by the fact values may be only partially ordered. We will elaborate on this point below.

In order to explain our framework, let us recall first the translation of LTL formulas to nondeterministic generalized Büchi automata (NGBW), as introduced in [30]. There, each state of the automaton is associated with a set of formulas, and the NGBW accepts a computation from a state q iff the computation satisfies exactly all the formulas associated with q . The state space of the NGBW contains only states associated with maximal and consistent sets of formulas, the transitions are defined so that requirements imposed by temporal formulas are satisfied, and the acceptance condition is used in order to guarantee that requirements that involve the satisfaction of eventualities are not delayed forever.

In the construction here, each state of the NGBW assigns a satisfaction value to every subformula. While it is not difficult to extend the local consistency rules to the latticed settings, handling of eventualities is more complicated. To see why, consider for example the formula Fp , for $p \in AP$, and the computation π in which the satisfaction value of p is $(\{a\}, \{b\}, \{c\})^\omega$. While $\llbracket \pi, Fp \rrbracket = \{a, b, c\}$, the computation never reaches a position in which the satisfaction value of the eventuality p is $\{a, b, c\}$. This poses a problem on translations of LTL formulas to automata, where eventualities are handled by making sure that each state in which the satisfaction of $\psi_1 U \psi_2$ is guaranteed is followed by a state in which the satisfaction of ψ_2 is guaranteed. For a multi-valued setting with fully-ordered values, as is the case in [12, 1], the latter can be replaced by a requirement to visit a state in which the guaranteed satisfaction value of ψ exceeds that of $\psi_1 U \psi_2$. As the example above demonstrates, such a position need not exist when the values are partially ordered. In order to address the above problem, every state in the NGBW associates with every subformula of the form $\psi_1 U \psi_2$ a value in \mathcal{L} that ψ_2 still needs “accumulate” in order for $\psi_1 U \psi_2$ to have its assigned satisfaction value. Thus, as in other break-point constructions [30, 22], we decompose the requirement to obtain a value ℓ to requirements to obtain join-irreducible values whose join is ℓ , and we check these requirements together.

Theorem 8. *Let φ be an LLTL formula over \mathcal{L} and $P \subseteq \mathcal{L}$ be a predicate. There exists an NGBW $\mathcal{A}_{\varphi, P}$ such that for every computation $\pi \in (2^{AP})^\omega$, it holds that $\llbracket \pi, \varphi \rrbracket \in P$ iff $\mathcal{A}_{\varphi, P}$ accepts π . The state space and transitions of $\mathcal{A}_{\varphi, P}$ are independent of P , which only influences the set of initial states. The NGBW $\mathcal{A}_{\varphi, P}$ has at most $|\mathcal{L}|^{O(|\varphi|)}$ states and index at most $|\varphi|$.*

Proof. We define $\mathcal{A}_{\varphi, P} = \langle \mathcal{L}^{AP}, Q, \delta, Q_0, \alpha \rangle$ as follows. Let $cl(\varphi)$ be the set of φ ’s subformulas, and let $ucl(\varphi)$ be the set of φ ’s subformulas of the form $\psi_1 U \psi_2$. Let G_φ and F_φ be the collection of functions $g : cl(\varphi) \rightarrow \mathcal{L}$ and $f : ucl(\varphi) \rightarrow \mathcal{L}$, respectively. For an element $v \in \mathcal{L}$, let $JI(v)$ be the minimal set $S \subseteq JI(\mathcal{L})$ such that $v = \bigvee_{s \in S} s$. By Birkhoff’s theorem, this set is well-defined, and the JI mapping is a bijection.

For a pair of functions $\langle g, f \rangle \in G_\varphi \times F_\varphi$, we say that $\langle g, f \rangle$ is *consistent* if for every $\psi \in cl(\varphi)$, the following holds.

- If $\psi = v \in \mathcal{L}$, then $g(\psi) = v$.
- If $\psi = \neg\psi_1$, then $g(\psi) = \neg g(\psi_1)$.
- If $\psi = \psi_1 \vee \psi_2$, then $g(\psi) = g(\psi_1) \vee g(\psi_2)$.

- If $\psi = \psi_1 \cup \psi_2$, then $\text{JI}(f(\psi)) \cap \text{JI}(g(\psi_2)) = \emptyset$.

The state space Q of $\mathcal{A}_{\varphi, \ell}$ is the set of all consistent pairs of functions in $G_\varphi \times F_\varphi$. Intuitively, while the function g describes the satisfaction value of the formulas in the closure, the function f describes, for each subformula of the form $\psi_1 \cup \psi_2$, the values in which ψ_2 still has to be satisfied in order for the satisfaction value $g(\psi_1 \cup \psi_2)$ to be fulfilled. Accordingly, if a value is in $\text{JI}(g(\psi_2))$, it can be removed from $f(\psi_1 \cup \psi_2)$, explaining why $\text{JI}(f(\psi_1 \cup \psi_2)) \cap \text{JI}(g(\psi_2)) = \emptyset$.

Then, $Q_0 = \{g \in Q : g(\varphi) \in P\}$ to contain all states in which the value assigned to φ is in P .

We now define the transition function δ . For two states $\langle g, f \rangle$ and $\langle g', f' \rangle$ in Q and a letter $\sigma \in \mathcal{L}^{AP}$, we have that $\langle g', f' \rangle \in \delta(\langle g, f \rangle, \sigma)$ iff the following hold.

- For all $p \in AP$, we have that $\sigma(p) = g(p)$.
- For all $X\psi_1 \in cl(\varphi)$, we have $g(X\psi_1) = g'(\psi_1)$.
- For all $\psi_1 \cup \psi_2 \in cl(\varphi)$, we have $g(\psi_1 \cup \psi_2) = g(\psi_2) \vee (g(\psi_1) \wedge g'(\psi_1 \cup \psi_2))$ and $f'(\psi_1 \cup \psi_2) = \begin{cases} \text{JI}(f(\psi_1 \cup \psi_2)) \setminus \text{JI}(g'(\psi_2)) & \text{If } \text{JI}(f(\psi_1 \cup \psi_2)) \neq \emptyset, \\ \text{JI}(g'(\psi_1 \cup \psi_2)) \setminus \text{JI}(g'(\psi_2)) & \text{Otherwise.} \end{cases}$

Finally, every formula $\psi_1 \cup \psi_2$ contributes to α the set $F_{\psi_1 \cup \psi_2} = \{\langle g, f \rangle : \text{JI}(f(\psi_1 \cup \psi_2)) = \emptyset\}$.

Observe that while δ is nondeterministic, it is only nondeterministic in the first component. That is, once the function g' is chosen, there is a single function f' that can match the transition. In Appendix A.2 we detail the correctness proof of the construction.

5 LLTL Synthesis

Recall that in the synthesis problem we are given an LLTL formula φ over sets I and O of input and output variables, taking truth values from a lattice \mathcal{L} , and we want to generate an (I/O) -transducer over \mathcal{L} all whose computations satisfy φ in a value from some desired set P of satisfaction values. In the noisy setting, the transducer may read a perturbed value of the input signals, and still all its computations need to satisfy φ as required. In this section we use the construction in Theorem 8 in order to solve both variants of the synthesis problem.

5.1 Solving the LLTL synthesis problem

Theorem 9. *The synthesis problem for LLTL is 2EXPTIME-complete. Given an LLTL formula φ over a lattice \mathcal{L} and a predicate $P \subseteq \mathcal{L}$, we can solve the synthesis problem for $\langle \varphi, P \rangle$ in time $2^{|\mathcal{L}|^{O(|\varphi|)}}$.*

Proof. Let m denote the size of \mathcal{L} , and let n denote the length of φ . The construction in Theorem 8 yields an NGBW with $m^{O(n)}$ and index n . By determinizing the NGBW we obtain an equivalent DPW $\mathcal{D}_{\varphi, P}$ with $2^{m^{O(n)} \log m^{O(n)}} = 2^{O(n)m^{O(n)}} = 2^{m^{O(n)}}$ states and index $m^{O(n)}$ [27, 23]. Following the same lines as the proof of Theorem 1, we see that in order to solve the LLTL synthesis problem, it suffices to solve the parity game that is obtained from \mathcal{D}_{φ} , except that here the alphabets of Players 1 and 2 are \mathcal{L}^O and \mathcal{L}^I . Accordingly, a winning memoryless strategy for Player 1 is an (I/O) -transducer over \mathcal{L} that realizes $\langle \varphi, P \rangle$.

Solving the parity game that is obtained from $\mathcal{D}_{\varphi, P}$ can be done (as we show in Theorem 1) in time $(2^{m^{O(n)}})^{m^{O(n)}} = 2^{m^{O(n)}}$. We conclude that the LLTL-synthesis problem is in 2EXPTIME. Hardness in 2EXPTIME follow from the hardness of the synthesis problem in the Boolean setting, which corresponds to a fully-ordered lattice with two values. \square

5.2 Solving the noisy LLTL synthesis problem

Consider an LLTL formula φ over the atomic propositions $I \cup O$, a predicate $P \subseteq \mathcal{L}$, and a noise function $\rho : \mathcal{L}^I \rightarrow 2^{\mathcal{L}^I}$. Recall that the goal in noisy synthesis is to find a transducer \mathcal{T} that realizes $\langle \varphi, P \rangle$ with noise ρ . Our goal is to construct a DPW on which we can apply the algorithm described in Theorem 1. For this, we proceed in three steps. First, we translate φ to a universal generalized co-Büchi word automaton (UGCW). Then, we incorporate the noise in the constructed UGCW. Finally, we determinize the UGCW to obtain a DPW, from which we proceed as described in Theorem 1. We start by showing how to incorporate noise in universal automata:

Lemma 1. *Consider a UGCW $\mathcal{D} = \langle I \cup O, Q, Q_0, \delta, \alpha \rangle$ and a noise function ρ . There exists a UGCW \mathcal{D}' such that \mathcal{D}' accepts a computation ρ iff \mathcal{D} accepts every computation ρ' that is ρ -indistinguishable from ρ . Moreover, \mathcal{D}' has the same state space and acceptance condition as \mathcal{D} .*

Proof. We obtain $\mathcal{D}' = \langle I \cup O, Q, Q_0, \delta', \alpha \rangle$ from \mathcal{D} by modifying δ as follows. For every $\sigma \in I \cup O$, let $\Gamma_\sigma = \{\gamma : \gamma|_O = \sigma|_O \text{ and } \gamma|_I \in \rho(\sigma|_I)\}$. Thus, Γ_σ contains all letters that are ρ -indistinguishable from σ . Then, for every state $q \in Q$, we have that $\delta'(q, \sigma) = \bigcup_{\gamma \in \Gamma_\sigma} \delta(q, \gamma)$. Thus, reading the letter σ , the UGCW \mathcal{D}' simulates all the runs of \mathcal{D} on all the letters that \mathcal{D} may read when the actual letter in the input is σ .

It is easy to show that the set of runs of \mathcal{D}' on a computation ρ is exactly the set of all the runs of \mathcal{D} on all the computations that are ρ -indistinguishable from ρ . From this, the correctness of the construction follows. \square

Theorem 10. *The noisy synthesis problem for LLTL is 2EXPTIME-complete. Given an LLTL formula φ over a lattice \mathcal{L} , a predicate $P \subseteq \mathcal{L}$, and a noise function ρ , we can solve the synthesis problem for $\langle \varphi, P \rangle$ with noise ρ in time $2^{m^{O(n)}}$.*

Proof. Let $\bar{P} = \mathcal{L} \setminus P$, and let $\mathcal{A}_{\varphi, \bar{P}}$ be the NGBW constructed for φ and \bar{P} in Theorem 8. Observe that $\mathcal{A}_{\varphi, \bar{P}}$ accepts a computation ρ iff $\llbracket \rho, \varphi \rrbracket \notin P$. Next, we dualize $\mathcal{A}_{\varphi, \bar{P}}$ and obtain a UGCW $\mathcal{D}_{\varphi, P}$ for the complement language, namely all computations ρ such that $\llbracket \rho, \varphi \rrbracket \in P$. We now apply the procedure in Lemma 1 to $\mathcal{D}_{\varphi, P}$ and obtain a UGCW $\mathcal{D}'_{\varphi, P}$ that accepts a computation ρ iff $\mathcal{D}_{\varphi, P}$ accepts every computation ρ' that is ρ -indistinguishable from ρ . Next, we determinize $\mathcal{D}'_{\varphi, P}$ to an equivalent DPW $\mathcal{D}''_{\varphi, P}$.

We claim that the algorithm described in the proof of Theorem 1 can be applied to $\mathcal{D}''_{\varphi, P}$. To see this, let $\mathcal{D}''_{\varphi, P} = \langle I \cup O, S, s_0, \eta, \beta \rangle$ and consider the game \mathcal{G} that is obtained from $\mathcal{D}''_{\varphi, P}$. That is, $\mathcal{G} = \langle \mathcal{L}^O, \mathcal{L}^I, S, s_0, \eta, \beta \rangle$, where for every $q \in S$, $i \in \mathcal{L}^I$, and $o \in \mathcal{L}^O$, we have that $\eta(q, i, o) = \mu(q, i \cup o)$.

A (memoryless) winning strategy f for Player 1 in \mathcal{G} is then an (I/O) -transducer over \mathcal{L} with the following property: for every strategy g of the environment, consider the play ρ that is induced by f and g . The play ρ then induces a computation $w \in \mathcal{L}^{I \cup O}$ that is accepted by $\mathcal{D}''_{\varphi, P}$. By the construction of $\mathcal{D}''_{\varphi, P}$, this means that for every computation w' that is ρ -indistinguishable from w , the run of $\mathcal{D}_{\varphi, P}$ on w' is accepting. Hence, $\llbracket w', \varphi \rrbracket \in P$, which in turn implies that f realizes $\langle \varphi, P \rangle$ with noise ρ .

Finally, we analyze the complexity of the algorithm. Let m denote the size of \mathcal{L} , and let n denote the length of φ . By Theorem 8, the size of $\mathcal{A}_{\varphi, \bar{P}}$ is $m^{O(n)}$ and it has index at most n . Dualizing results in a UGCW of the same size and acceptance condition, and so is the transition to $\mathcal{D}'_{\varphi, P}$. Determinization involves an exponential blowup, such that $\mathcal{D}''_{\varphi, P}$ has $2^{m^{O(n)} \log m^{O(n)}} = 2^{m^{O(n)}}$ states and index $m^{O(n)}$. Finally, solving the parity game can be

done in time $(2^{m^{O(n)}})^{m^{O(n)}} = 2^{m^{O(n)}}$. We conclude that the LLTL-noisy-synthesis problem is in 2EXPTIME. Hardness in 2EXPTIME again follows from the hardness of the synthesis problem in the Boolean setting. \square

Remark 2. The approach described in the proofs of Theorems 1, 9, and 10 is Safrafull, in the sense it involves a construction of a DPW. As has been the case with Boolean synthesis [20], it is possible to proceed Safralessly also in LLTL synthesis with noise. To see this, note that the starting point in Theorem 1 can also be a UGCW, and that Lemma 1 works with UGCWs. In more details, once we construct a UGCW \mathcal{U} for the specification, possibly with noise incorporated, the Safraless approach expands \mathcal{U} to a universal co-Büchi tree automaton that accepts winning strategies for the system in the corresponding synthesis game, and checks its emptiness. In terms of complexity, rather than paying an additional exponent in the translation of the specification to a deterministic automaton, we pay it in the non-emptiness check of the tree automaton.

5.3 Local stability revisited

In Section 3.2 we have seen that not all LLTL formulas are locally stable. This gives rise to the question of deciding whether a given LLTL formula is locally stable. In the context of synthesis, if φ is known to be locally stable and we have a transducer \mathcal{T} that realizes $\langle \varphi, P \rangle$ with no noise, we know that \mathcal{T} realizes $\langle \varphi, P \oplus \gamma \rangle$ with noise ρ_γ , where $\rho_\gamma(\sigma) = \{\tau : d(\sigma, \tau) \leq \gamma\}$, and $P \oplus \gamma$ is the extension of P to noise ρ_γ . Thus, $\ell \in P \oplus \gamma$ iff there is $\ell' \in P$ such that $d(\ell, \ell') \leq \gamma$.

More generally, we consider the following problem: given an LLTL formula φ and a noise-threshold γ , we want to compute the maximal *distraction*, denoted $\Delta_{\varphi, \gamma}$, that noise γ may cause to φ . Formally,

$$\Delta_{\varphi, \gamma} = \max \{d(\llbracket \pi, \varphi \rrbracket, \llbracket \pi', \varphi \rrbracket) : \pi, \pi' \in (\mathcal{L}^{AP})^\omega \text{ and } ld(\pi, \pi') \leq \gamma\}.$$

Observe that finding $\Delta_{\varphi, \gamma}$ allows us to decide local stability by iterating over all elements $\gamma \in \{1, \dots, |\mathcal{L}|\}$ and verifying that $\Delta_{\varphi, \gamma} \leq \gamma$. To simplify things, we work with a threshold version of the problem:

Lemma 2. *Given an LLTL formula φ and thresholds γ and μ , the problem of deciding whether $\Delta_{\varphi, \gamma} \leq \mu$ is in PSPACE.*

Proof. We solve the dual problem, namely deciding whether there exist $\pi, \pi' \in (\mathcal{L}^{AP})^\omega$ such that $ld(\pi, \pi') \leq \gamma$ and $d(\llbracket \pi, \varphi \rrbracket, \llbracket \pi', \varphi \rrbracket) > \mu$. In order to solve this problem, we proceed as follows. In Theorem 8 we showed how to construct a NGBW $\mathcal{A}_{\varphi, \ell}$ such that $\mathcal{A}_{\varphi, \ell}$ accepts a computation π iff $\llbracket \pi, \varphi \rrbracket = \ell$. In Section 5.2, we showed how to construct a UGCW $\mathcal{D}'_{\varphi, \ell \oplus \mu}$ such that $\mathcal{D}'_{\varphi, \ell \oplus \mu}$ accepts π iff $\llbracket \pi', \varphi \rrbracket \in \ell \oplus \mu$ for every computation π' that is ρ_γ -indistinguishable from π . Now, there exist $\pi, \pi' \in (\mathcal{L}^{AP})^\omega$ such that $ld(\pi, \pi') \leq \gamma$ and $d(\llbracket \pi, \varphi \rrbracket, \llbracket \pi', \varphi \rrbracket) > \mu$ iff there exists $\ell \in \mathcal{L}$ such that $\llbracket \pi, \varphi \rrbracket = \ell$ and the latter conditions hold. Observe that these conditions hold iff there exists a computation π that is accepted by $\mathcal{A}_{\varphi, \ell}$ but not by $\mathcal{D}'_{\varphi, \ell \oplus \mu}$. Thus, it suffices to decide whether $L(\mathcal{A}_{\varphi, \ell}) \cap \overline{L(\mathcal{D}'_{\varphi, \ell \oplus \mu})} = \emptyset$ for every $\ell \in \mathcal{L}$.

Finally, we analyze the complexity of this procedure. Let $|\mathcal{L}| = m$ and $|\varphi| = n$. Complementing of $\mathcal{D}'_{\varphi, \ell \oplus \mu}$ can be done by constructing $\mathcal{D}'_{\varphi, \ell \oplus \mu}$. Hence, both $\mathcal{A}_{\varphi, \ell}$ and $\overline{\mathcal{D}'_{\varphi, \ell \oplus \mu}}$ have $m^{O(n)}$ states. Checking the emptiness of their intersection can be done on-the-fly in PSPACE. \square

As discussed above, Lemma 2 suggest a PSPACE algorithm for deciding local stability. We now complete the picture by presenting a matching lower bound.

Theorem 11. *Given an LLTL formula φ over a lattice \mathcal{L} , deciding whether φ is locally stable is PSPACE-complete.*

Proof. Membership in PSPACE follows from Lemma 2. We prove hardness by describing a polynomial time reduction from the satisfiability problem for LTL to the complement of the local-stability problem.

Consider an LTL formula φ over AP . We assume that φ is not valid, thus there is a computation that does not satisfy it (clearly LTL satisfiability is PSPACE-hard also with this promise). We construct an LLTL formula ψ over the lattice $\mathcal{L} = \langle 2^{\{a,b\}}, \subseteq \rangle$ as follows. Let $AP' = \{p' : p \in AP\}$ be a tagged copy of AP . We define $\psi = \varphi \vee \varphi'$ over $AP \cup AP'$, where φ' is obtained from φ by replacing each atomic proposition by its tagged copy. Clearly this reduction is polynomial. We claim that φ is satisfiable iff ψ is not locally stable.

Consider φ as an LLTL formula over \mathcal{L} . We observe that if there exists a computation π such that $\llbracket \pi, \varphi \rrbracket = \{a\}$, then there exists a computation τ such that $\llbracket \tau, \varphi \rrbracket = \{b\}$. Indeed, the lattice \mathcal{L} is symmetric and φ does not contain elements of the form $\{a\}$ or $\{b\}$ to break the symmetry. Thus, we can obtain τ by swapping the roles of a and b in π . From Theorem 3 we know that \mathcal{L} is pointed, so in this case there also exists a computation ρ such that $\llbracket \rho, \varphi \rrbracket = \{a, b\}$.

We first prove that if φ is not satisfiable, then ψ is locally stable. Observe that if we view φ as an LLTL formula and there exists a computation π such that $\llbracket \pi, \varphi \rrbracket = \{a\}$, then φ is satisfiable as an LTL formula. Indeed, a satisfying computation π' for φ can be obtained from π by defining $p \in \pi'_i$ iff $a \in \pi_i(p)$ for all $p \in AP$ and $i \geq 0$. Thus, if φ is not satisfiable, then $\llbracket \pi, \varphi \rrbracket = \emptyset$ for every computation $\pi \in (\mathcal{L}^{AP})^\omega$, and similarly $\llbracket \pi', \psi \rrbracket = \emptyset$ for every $\pi' \in (\mathcal{L}^{AP \cup AP'})^\omega$. Hence, ψ is locally stable.

For the second direction, assume that φ is satisfiable. Thus, there exists a computation $\pi \in (2^{AP})^\omega$ such that $\pi \models \varphi$. By our assumption, there also exists a computation π' such that $\pi' \not\models \varphi$. It is easy to see that by identifying True with $\{a, b\}$ and False with \emptyset , we get $\llbracket \pi, \varphi \rrbracket = \{a, b\}$ and $\llbracket \pi', \varphi \rrbracket = \emptyset$. For a computation $w \in (\mathcal{L}^{AP})^\omega$, let $\hat{w} \in (\mathcal{L}^{AP \cup AP'})^\omega$ be the computation obtained from w by copying the behavior of the atoms in AP to their tagged atoms. Thus, for all $i \geq 0$, we have $p, p' \in \hat{w}_i$ iff $p \in w_i$.

Since the maximal distance between elements in \mathcal{L} is 2, there exists a computation τ such that $ld(\pi', \tau) \leq 1$ and $ld(\tau, \pi) \leq 1$. That is, we can “get” from π' to π by two local changes of 1. From this follows that $ld(\hat{\pi}', \hat{\tau}) \leq 1$ and $ld(\hat{\tau}, \hat{\pi}) \leq 1$. Consider $\llbracket \tau, \varphi \rrbracket$. If $\llbracket \tau, \varphi \rrbracket = \{a, b\}$, then ψ is not locally stable, since $ld(\hat{\pi}', \hat{\tau}) \leq 1$ but $d(\llbracket \hat{\tau}, \psi \rrbracket, \llbracket \hat{\pi}', \psi \rrbracket) = 2$. Similarly, if $\llbracket \tau, \varphi \rrbracket = \emptyset$, then ψ is not locally stable. Otherwise, w.l.o.g. $\llbracket \tau, \varphi \rrbracket = \{a\}$. Then, there exists a computation τ' such that $\llbracket \tau', \varphi \rrbracket = \{b\}$. The computation τ' is obtained by swapping a and b in τ . Observe that since π' contains only symmetric elements from \mathcal{L} (i.e. \emptyset and $\{a, b\}$), then $ld(\pi', \tau) = ld(\pi', \tau') \leq 1$. Finally, since AP and AP' are disjoint, then by assigning τ over AP and τ' over AP' , we obtain a computation ρ such that $ld(\hat{\pi}', \rho) \leq 1$ but $\llbracket \rho, \psi \rrbracket = \{a\} \vee \{b\} = \{a, b\}$, and ψ is not locally stable. \square

References

1. S. Almagor, U. Boker, and O. Kupferman. Formalizing and reasoning about quality. In *Proc. 40th ICALP, LNCS 7966*, pages 15–27, 2013.
2. R. Alur, A. Kanade, and G. Weiss. Ranking automata and games for prioritized requirements. In *Proc. 20th CAV, LNCS 5123*, pages 240–253, 2008.
3. R. Bloem, K. Chatterjee, T. Henzinger, and B. Jobstmann. Better quality in synthesis through quantitative objectives. In *Proc. 21st CAV, LNCS 5643*, pages 140–156, 2009.

4. P. Cerný and T. Henzinger. From boolean to quantitative synthesis. In *EMSOFT*, pages 149–154, 2011.
5. K. Chatterjee, L. Doyen, and T. Henzinger. Quantative languages. In *Proc. 17th CSL*, pages 385–400, 2008.
6. K. Chatterjee, L. Doyen, T. A. Henzinger, and J.-F. Raskin. Algorithms for omega-regular games with imperfect information. In *Proc. 15th CSL, LNCS 4207*, pages 287–302, 2006.
7. K. Chatterjee and R. Majumdar. Minimum attention controller synthesis for omega-regular objectives. In *FORMATS*, pages 145–159, 2011.
8. K. Chatterjee, R. Majumdar, and T. A. Henzinger. Controller synthesis with budget constraints. In *Proc 11th HSCC, LNCS 4981*, pages 72–86, 2008.
9. M. Chechik, B. Devereux, and A. Gurfinkel. Model-checking infinite state-space systems with fine-grained abstractions using SPIN. In *Proc. 8th Int. SPIN Workshop, LNCS 2057*, pages 16–36, 2001.
10. A. Church. Logic, arithmetics, and automata. In *Proc. Int. Congress of Mathematicians, 1962*, pages 23–35. Institut Mittag-Leffler, 1963.
11. E. Emerson and C. Jutla. Tree automata, μ -calculus and determinacy. In *Proc. 32nd FOCS*, pages 368–377, 1991.
12. M. Faella, A. Legay, and M. Stoelinga. Model checking quantitative linear time logic. *Electr. Notes Theor. Comput. Sci.*, 220(3):61–77, 2008.
13. E. Filiot, N. Jin, and J.-F. Raskin. An antichain algorithm for LTL realizability. In *Proc. 21st CAV, LNCS 5643*, pages 263–277, 2009.
14. M. Huth and S. Pradhan. Consistent partial model checking. *Electr. Notes Theor. Comput. Sci.*, 73:45–85, 2004.
15. M. Jurdzinski, M. Paterson, and U. Zwick. A deterministic subexponential algorithm for solving parity games. *SIAM Journal on Computing*, 38(4):1519–1532, 2008.
16. R. Kumar and M. Shayman. Supervisory control of nondeterministic systems under partial observation and decentralization. *SIAM Journal of Control and Optimization*, 1995.
17. O. Kupferman and Y. Lustig. Lattice automata. In *Proc. 8th VMCAI, LNCS 4349*, pages 199 – 213, 2007.
18. O. Kupferman and Y. Lustig. Latticed simulation relations and games. *International Journal on the Foundations of Computer Science*, 21(2):167–189, 2010.
19. O. Kupferman and M. Vardi. Church’s problem revisited. *The Bulletin of Symbolic Logic*, 5(2):245 – 263, 1999.
20. O. Kupferman and M. Vardi. Safrless decision procedures. In *Proc. 46th FOCS*, pages 531–540, 2005.
21. M. Kwiatkowska. Quantitative verification: models techniques and tools. In *ESEC/SIGSOFT FSE*, pages 449–458, 2007.
22. S. Miyano and T. Hayashi. Alternating finite automata on ω -words. *TCS*, 32:321–330, 1984.
23. N. Piterman. From nondeterministic Büchi and Streett automata to deterministic parity automata. In *Proc. 21st LICS*, pages 255–264. IEEE press, 2006.
24. A. Pnueli and R. Rosner. On the synthesis of a reactive module. In *Proc. 16th POPL*, pages 179–190, 1989.
25. A. Pnueli and R. Rosner. On the synthesis of an asynchronous reactive module. In *Proc. 16th ICALP, LNCS 372*, pages 652–671. Springer, 1989.
26. J. Reif. The complexity of two-player games of incomplete information. *Journal of Computer and Systems Science*, 29:274–301, 1984.
27. S. Safra. Exponential Determinization for ω -Automata with Strong-Fairness Acceptance Condition. In *Proc. 24th STOC*, pages 275–282, 1992.
28. S. Schewe. Solving Parity Games in Big Steps. In *Proc. 27th FSTTCS*, pages 449–460, 2007.
29. M. Vardi. From verification to synthesis. In *Proc. 2nd Int. Conf. on Verified Software: Theories, Tools, Experiments, LNCS 5295*, page 2, 2008.
30. M. Vardi and P. Wolper. Reasoning about infinite computations. *Information and Computation*, 115(1):1–37, 1994.
31. Y. Velner and A. Rabinovich. Church synthesis problem for noisy input. In *Proc. 14th FOSSACS*, pages 275–289, 2011.

A Proofs

A.1 Completing the proof of Theorem 5

To simplify the proof, we observe that since π and π' differ only in a single index, and in particular only in a finite prefix, we can avoid treating the case of \cup subformulas. Indeed, we can expand subformulas of the form $\psi \cup \theta$ using propositional conjunctives and the X operator so that all the suffixes before π^i are not evaluated on subformulas that contain \cup . Clearly, the value of the remaining evaluation of \cup subformulas does not change, as π and π' agree on suffixes that start after the i -th position.

- If $\varphi = \ell \in \mathcal{L}$ the claim is trivial.
- If $\varphi = p \in AP$, then $d(\llbracket \pi, \varphi \rrbracket, \llbracket \pi', \varphi \rrbracket) = d(\llbracket \pi_0, p \rrbracket, \llbracket \pi'_0, p \rrbracket)$, which, by the assumption, is at most 1. Moreover, $\llbracket \pi', p \rrbracket \in \{\llbracket \pi, p \rrbracket, \llbracket \pi, p \rrbracket \vee u\}$.
- If $\varphi = \neg\psi$, then by the induction hypothesis it holds that $\llbracket \pi', \psi \rrbracket \in \{\llbracket \pi, \psi \rrbracket \wedge \neg u, \llbracket \pi, \psi \rrbracket, \llbracket \pi, \psi \rrbracket \vee u\}$ and $d(\llbracket \pi, \psi \rrbracket, \llbracket \pi', \psi \rrbracket) \leq 1$. Hence, by the properties of negation, $d(\neg\llbracket \pi, \psi \rrbracket, \neg\llbracket \pi', \psi \rrbracket) \leq 1$ and $\llbracket \pi', \varphi \rrbracket \in \{\llbracket \pi, \varphi \rrbracket \wedge \neg u, \llbracket \pi, \varphi \rrbracket, \llbracket \pi, \varphi \rrbracket \vee u\}$, and we are done.
- If $\varphi = \psi \vee \theta$, then, by the induction hypothesis, it holds in particular that $\llbracket \pi, \psi \rrbracket \wedge \neg u \leq \llbracket \pi', \psi \rrbracket \leq \llbracket \pi, \psi \rrbracket \vee u$ and $\llbracket \pi, \theta \rrbracket \wedge \neg u \leq \llbracket \pi', \theta \rrbracket \leq \llbracket \pi, \theta \rrbracket \vee u$. Therefore, $\llbracket \pi, \varphi \rrbracket \wedge \neg u \leq \llbracket \pi', \varphi \rrbracket \leq \llbracket \pi, \varphi \rrbracket \vee u$. Figure 3 demonstrates the relations. Intuitively, one can see that any additional lattice element between $\psi \vee \theta \vee u$ and $(\psi \vee \theta) \wedge \neg u$ induces an N5-structure. Thus, $\llbracket \pi', \varphi \rrbracket$ must be within distance 1 of $\llbracket \pi, \varphi \rrbracket$. This is the key point in the proof.

Formally, in order to prove that the distance is preserved, we distinguish between cases. First, if $\llbracket \pi', \varphi \rrbracket = \llbracket \pi, \varphi \rrbracket$, then we are done. If $\llbracket \pi', \varphi \rrbracket = \llbracket \pi, \varphi \rrbracket \vee v$, then since $\llbracket \pi', \varphi \rrbracket = \llbracket \pi', \psi \rrbracket \vee \llbracket \pi', \theta \rrbracket$, it must be that w.l.o.g $\llbracket \pi', \psi \rrbracket = \llbracket \pi, \psi \rrbracket \vee u$. Assume by way of contradiction that there exists $t \in \mathcal{L}$ such that $\llbracket \pi', \varphi \rrbracket < t < \llbracket \pi, \varphi \rrbracket \vee v$. We then have the N5 structure³ $\langle \psi \vee \theta, t, \psi \vee u, \psi, \psi \vee \theta \vee u \rangle$. Since, however, \mathcal{L} is distributive, it cannot have an N5 structure, we have reached a contradiction.

The case $\llbracket \pi, \varphi \rrbracket \wedge \neg u \leq \llbracket \pi', \varphi \rrbracket$ is handled similarly.




Fig. 3. Theorem 5 - disjunction case.

³ It may be the case that some of the nodes coincide, and this is not a proper N5. However, these cases are easy to handle.

– If $\varphi = X\psi$, the claim follows immediately from the induction hypothesis. \square

A.2 Completing the proof of Theorem 8

We now proceed to prove the correctness of the construction and analyze its complexity. In the proof we identify a set $S \subseteq \text{JI}(\mathcal{L})$ with the element $\bigvee_{s \in S} s \in \mathcal{L}$. Observe that it suffices to prove that for every $\ell \in \mathcal{L}$, the NGBW $\mathcal{A}_{\varphi, \{\ell\}}$ accepts a computation π iff $\llbracket \pi, \varphi \rrbracket = \ell$. We first prove that if $\pi \in (\mathcal{L}^{AP})^\omega$ is such that $\llbracket \pi, \varphi \rrbracket = \ell$ for some $\ell \in \mathcal{L}$, then $\mathcal{A}_{\varphi, \{\ell\}}$ accepts π . For every $i \in \mathbb{N}$, let $g_i \in G_\varphi$ be such that for all $\psi \in cl(\varphi)$ we have that $g_i(\psi) = \llbracket \pi^i, \psi \rrbracket$. Also, let $f_0 : ucl(\varphi) \rightarrow \mathcal{L}$ be such that for every subformula of the form $\psi_1 U \psi_2$ we have $f_0(\psi_1 U \psi_2) = \text{JI}(g_0(\psi_1 U \psi_2)) \setminus \text{JI}(g_0(\psi_2))$. Finally, for $i \in \mathbb{N}$, let f_{i+1} be induced from f_i and g_{i+1} in the single way that satisfies the conditions in the definition of δ .

We claim that $r = \langle g_0, f_0 \rangle, \langle g_1, f_1 \rangle, \dots$ is an accepting run of $\mathcal{A}_{\varphi, \{\ell\}}$ on π . First, the semantics of LLTL implies that the consistency conditions, both the local ones and these imposed by δ are satisfied. In particular, for the conditions imposed by δ , this follows from the fact that for all positions $i \in \mathbb{N}$ we have that $\llbracket \pi^i, X\psi_1 \rrbracket = \llbracket \pi^{i+1}, \psi_1 \rrbracket$ and $\llbracket \pi^i, \psi_1 U \psi_2 \rrbracket = \llbracket \psi_2 \rrbracket \vee (\llbracket \pi^i, \psi_1 \rrbracket \wedge \llbracket \pi^i, \psi_1 U \psi_2 \rrbracket)$. Also, since $g_0(\varphi) = \ell$, then $\langle g_0, f_0 \rangle \in Q_0$.

It is left to prove that r is accepting. Consider a sub-formula of the form $\psi_1 U \psi_2$. We prove that r visits $F_{\psi_1 U \psi_2}$ infinitely often. Consider a position $i \in \mathbb{N}$ and let $\llbracket \pi^i, \psi_1 U \psi_2 \rrbracket = y$. We prove that there is a position $n \geq i$ such that $f_n = \emptyset$, thus $\langle g_n, f_n \rangle \in F_{\psi_1 U \psi_2}$. By the semantics of U and the finiteness of \mathcal{L} , there is a (minimal) index $n \geq i$ such that $y = \bigvee_{i \leq j \leq n} (\llbracket \pi^j, \psi_2 \rrbracket \wedge \bigwedge_{i \leq k < j} \llbracket \pi^k, \psi_1 \rrbracket)$. It is easy to prove by induction on $n - i$ that there exists some $i \leq k \leq n$ such that $\text{JI}(f_k(\psi_1 U \psi_2)) = \emptyset$, using the fact that $f_j(\psi_1 U \psi_2) \leq g_j(\psi_1 U \psi_2)$ for all $j \geq 0$.

The other direction is more complicated. Let $\pi \in (\mathcal{L}^{AP})^\omega$ be such that π is accepted by $\mathcal{A}_{\varphi, \{\ell\}}$. We prove that $\llbracket \pi, \varphi \rrbracket = \ell$. Let $\rho = \langle g_1, f_1 \rangle, \langle g_2, f_2 \rangle, \dots$ be an accepting run of $\mathcal{A}_{\varphi, \{\ell\}}$ on π , and let $h_1, h_2, \dots \in (G_\varphi)^\omega$ be such that for all $i \in \mathbb{N}$ and $\psi \in cl(\varphi)$ we have that $h_i(\psi) = \llbracket \pi^i, \psi \rrbracket$. We claim that $h_i = g_i$ for all $i \in \mathbb{N}$. The proof is by induction on the structure of the formulas in $cl(\varphi)$. Consider a formula $\psi \in cl(\varphi)$. If $\psi = p \in AP$, then since ρ is a legal run, a transition from state $\langle g_i, f_i \rangle$ is possible with letter σ iff $\sigma(p) = \llbracket \pi^i, p \rrbracket = \rho(p)$, and we are done. If $\psi = v \in \mathcal{L}$, $\psi = \psi_1 \vee \psi_2$, or $\psi = X\psi_1$, then the claim follows from the consistency rules and the induction hypothesis. Finally, if $\psi = \psi_1 U \psi_2$, then, as we prove in Lemma 3 below, the fact ρ is an accepting run implies the first equality in the chain below. The second equality follows from the induction hypothesis, and the third equality is from the semantics of LLTL.

$$g_i(\psi) = \bigvee_{i \leq j} (g_j(\psi_2) \wedge \bigwedge_{i \leq k < j} g_k(\psi_1)) = \bigvee_{i \leq j} (\llbracket \pi^j, \psi_2 \rrbracket \wedge \bigwedge_{i \leq k < j} \llbracket \pi^k, \psi_1 \rrbracket) = \llbracket \pi^i, \psi \rrbracket.$$

We conclude that $h_0 = g_0$. Since $g_0 \in Q_0$, it follows that $\llbracket \pi, \varphi \rrbracket = \ell$ and we are done. \square

Lemma 3. $g_i(\psi_1 U \psi_2) = \bigvee_{i \leq j} (g_j(\psi_2) \wedge \bigwedge_{i \leq k < j} g_k(\psi_1))$.

Proof. Since ρ is a legal run, then for every $i \in \mathbb{N}$, it holds that

$$g_i(\psi_1 U \psi_2) = g_i(\psi_2) \vee (g_i(\psi_1) \wedge g_{i+1}(\psi_1 U \psi_2)). \quad (*)$$

We prove the lemma by proving that for every $v \in \text{JI}(\mathcal{L})$ it holds that $g_i(\psi_1 U \psi_2) \geq v$ iff $\bigvee_{i \leq j} (g_j(\psi_2) \wedge \bigwedge_{i \leq k < j} g_k(\psi_1)) \geq v$. The equality then follows from Birkhoff's theorem.

Using (*), it is easy to prove by induction that for every index i and for every $n \in \mathbb{N}$ it holds that

$$g(\psi_1 \cup \psi_2) = \bigvee_{i \leq j \leq n} \left(g_j(\psi_2) \wedge \bigwedge_{i \leq k < j} g_k(\psi_1) \right) \vee \left(g_{n+1}(\psi_1 \cup \psi_2) \wedge \bigwedge_{i \leq k \leq n} g_k(\psi_1) \right) \quad (**)$$

Let $v \in \text{JI}(\mathcal{L})$ and assume that

$$\bigvee_{i \leq j} (g_j(\psi_2) \wedge \bigwedge_{i \leq k < j} g_k(\psi_1)) \geq v.$$

Since v is join-irreducible, it follows that there exist some $n \in \mathbb{N}$ such that

$$g_n(\psi_2) \wedge \bigwedge_{i \leq k < n} g_k(\psi_1) \geq v.$$

Since (**) is true for every n , then in particular we have that $g(\psi_1 \cup \psi_2) \geq v$, which concludes the first direction of the proof.

For the second direction, assume that $g(\psi_1 \cup \psi_2) \geq v$. If there exists $n \geq i$ such that $\bigvee_{i \leq j \leq n} (g_j(\psi_2) \wedge \bigwedge_{i \leq k < j} g_k(\psi_1)) \geq v$, then we are done. Assume by way of contradiction that there is no such n . Thus, by (**), for every $n \geq i$ it holds that $g_{n+1}(\psi_1 \cup \psi_2) \wedge \bigwedge_{i \leq k \leq n} g_k(\psi_1) \geq v$, which means that $g_{n+1}(\psi_1 \cup \psi_2) \geq v$ and $\bigwedge_{i \leq k \leq n} g_k(\psi_1) \geq v$. In particular, there cannot exist $n \geq i$ such that $g_n(\psi_2) \geq v$, otherwise it would contradict our assumption.

Since the run is accepting, there exists $n_1 > i$ such that $f_{n_1}(\psi_1 \cup \psi_2) = \emptyset$. Consider the suffix of the run starting from $n_1 + 1$. For every $t \geq n_1$, We have that $g_t(\psi_1 \cup \psi_2) \geq v$ but $g_t(\psi_2) \not\geq v$. Thus, $v \in \text{JI}(f_{n_1+1}(\psi_1 \cup \psi_2))$, and v will never be removed from the f component, this is in contradiction to the fact that the run is accepting, and we are done. \square